Model Checking concurrent objects

Submit Time: Trinity Term, 2022

Degree: Part B, Computer Science

Candidate Number: 1043026

Word count(excluding figure): 9167

Abstract

Concurrent objects are convenient tools for programmers. With concurrent objects, programmers can write code with multiple threads as if they are writing a single-threaded code. However, it is crucial to know the correctness of concurrent objects. In this thesis, we study a technique to justify the correctness of synchronization objects. We first present CSP models for common concurrent primitives according to their behaviors. We then systematically build several concurrent objects from their Scala sources. We make assertions of these synchronization objects with a technique derived from linearizability testing. We find these assertions can effectively find bugs in a concurrent datatype and provide a history to give more context for the developer.

Contents

1 Introduction

Concurrent objects are convenient tools for programmers. With concurrent datatypes, programmers can write code with multiple threads as if they are writing a single-threaded code. However, it is crucial to know the correctness of concurrent objects. If the implementation of a concurrent object is wrong, then code using the concurrent object is very likely to be faulty.

The Channel object is one synchronization object commonly used in Go. The channel object can be used to share data from one process to another process. A process can send data to another process by calling send with the data to share. Likewise, a process can receive data from other processes by calling the receive function. In Go and Communicating Scala Object package, channels are unbuffered by default. If there is no process to receive the data, the sending process blocks until a process is willing to receive its data. Similarly, a receiving process blocks until a process sends some data.

```
trait Channel{
  def send(data: Int): Int
  def receive(): Int
}
```

Figure 1: Interface of a Channel object

In this thesis, we shall study the correctness of synchronization objects. Each synchronization in a synchronization object involves multiple processes, whereas synchronization in concurrent datatypes like concurrent queue and concurrent only involves a single process.

There are two main properties to check for a synchronization object, the safety property and the liveness property. The safety property states that the history of the synchronization object should satisfy some conditions. For example, if one process sends 1 when no other process is sending, then a process calling receive should only receive1. The liveness property states that the concurrent object should not refuse to synchronize when synchronization is possible between one or more processes. For

example, if a process calls send and a process receive, the system should be able to synchronize and should not deadlock.

1.1 Thesis Overview

In the remaining part of Section 1, we describe the correctness condition for a synchronization object and abstractly how to test these conditions in CSP using the linearization test technique.

In section 2, we build CSP modules for common concurrent primitives such as shared variables, monitors and semaphores.

Starting from section 3, we use the linearization test technique to distinguish correct and faulty implementation for several synchronization objects from [Low21]. We first implement the synchronization object in CSP according to its Scala source code. Then we write specifications for a system using the synchronization object and carry out the tests.

1.2 Synchronization linearizability test

To verify the correctness of a concurrent datatype, one can carry out the linearizability test described in the paper Testing for Linearizability [Low17]. The linearizability testing framework logs the orders of each function call and function return. The testing framework then justifies the observed history by attempting to find a series of synchronization points that obey the safety property. The concurrent datatype implementation is considered faulty if the framework can not find a valid synchronization point series.

In this remaining section, we shall look at a few examples of histories of systems using the Channel object. Figure 2 visualizes the history of a system with two processes. Process T1 calls send with argument 1 and returns. Process T2 calls receive and returns with 1. Each long horizontal line in the timeline represents a function call made by the corresponding process. The short vertical bars at the two ends of

the long horizontal line indicate the function call's starting time and ending time. And the long vertical line between T1 and T2 represents the synchronization between the two processes.

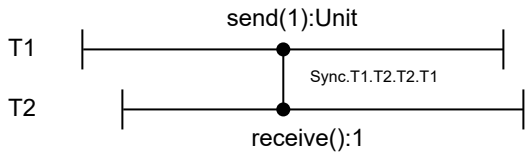

Figure 2: Visualized history of T1 calling send(1) and T2 calling receive()

Figure 3 shows a timeline similar to Figure 2, but T2 returns 2 instead of 1. In this case, the linearizability test framework can not justify the return of processT2's receive, and suggests the trace is generated by a faulty channel implementation.

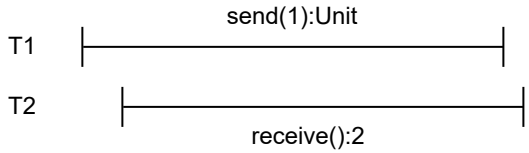

Figure 3: Visualized history of T1 sends 1 but T2 receives 2

In Figure 4, both processes calls send, and no synchronization is possible. Note that the liveness condition is not invalidated even if the system deadlocks in this case.

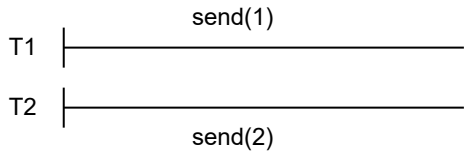

Figure 4: Visualized history of both T1 and T2 calling send

Scheduling is one of the reasons validating a history can be complicated. In Figure 5, process T3 calls send(3) first but gets descheduled. Then T1 calls send(1) and synchronizes with T2 which later calls receive. The linearization framework

usually needs to search a large state space to find a valid series of synchronization points.

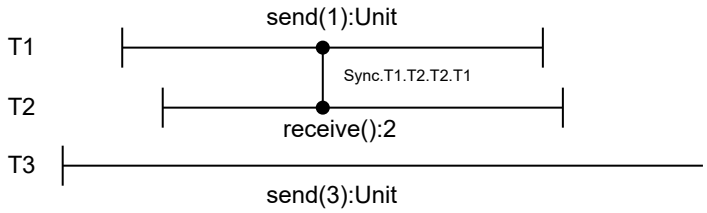

Figure 5: Visualized history of T3 get descheduled

1.3 Checking safety property using CSP

The history observed by the linearizability framework can be captured as a trace in CSP. A Call event in CSP represents the start of a function call in the observed history. A Return event represents the returning of a function call. For the safety property, we check that set of all possible histories of a testing system is a subset of all correct histories. In CSP, this corresponds to an assertion that a testing system trace refines a specification of systems using the synchronization object. And we can use FDR [GRABR14] to check the assertion.

A generic and scalable system is used as the testing system. Each process in the testing system can call any function from the concurrent object with any arguments allowed. Each process must be allowed to terminate. Otherwise, the testing system only models a system that runs forever, given that there is no deadlock. We shall see how this affects bug finding in a concurrent datatype in later objects.

The specification process is constructed using the linearization technique. On the high level, the specification process for the system internally uses Sync events to represent synchronization between processes. Inside the specification process, some sub-processes generate corresponding Call and Return event for every synchronization point. When all sub-processes are placed in parallel, the Sync event agrees. So the resulting specification system generates all possible histories.

We shall see a concrete implementation of a testing system and a specification process in the MenWomen section.

1.4 Checking liveness property using CSP

For liveness property, we check the same generic and scalable testing system refines the same specification process, but in the failure model. One could use a datatype-specific specification process that does not explicitly use any synchronization points. However, reusing the linearizer process is easier.

1.5 Related work

Testing for Linearizability [Low17] describes a framework to test concurrent datatypes. The testing framework uses observations of histories. As the observed system does not generate traces in a specific order, the framework may use a long time to exhaust all histories. Also, the justifying algorithm cannot check infinite traces. However, the linearization test can check all finite and infinite histories of a system using FDR's state machine.

There are also runtime programming tools to detect race conditions and deadlocks in concurrent code. Thread Sanitizer [SI09] detects race conditions and deadlocks in C++ and Go when the program is running. Runtime checkers often bring extra overheads, which may not represent the real-world scenario. Also, a test with runtime programming tools is neither unlikely to exhaust all possible histories efficiently.

In Chapter 19 of Understanding Concurrent Systems [Ros10], the author describes a CSP model for shared variables and provides a tool to analyze shared variable programs. This project further supports analyzing programs using more synchronizations, such as monitors and semaphores.

2 Common Objects

2.1 Shared Variable

The usage of shared variables is common in concurrent datatypes. For example, some concurrent datatypes may temporarily store the identity of a waiting process. However, CSP is more like a functional programming language and does not support mutable variables.

A recursive process in CSP can capture the behavior of a shared variable. The recursive process holds the value of the variable in its parameter. At any time, the variable process is willing to answer a query for the variable value in channel getValue. Alternatively, the process can receive an update on the variable value in channel getValue, after which the function recurses with the new variable value.

Because it is natural for a concurrent datatype to use multiple shared variables, the global variable is implemented as a CSP module in Figure 6 to allow better code reuse. The module requires two parameters. TypeValue is the set of possible values for the variable, and initialValue is the value before any process modifies the variable. An uninitialized variable module is also available in the same Figure 6, with the only difference that the variable non-deterministically chooses an initial value from TypeValue at start time. runWith is a convenient helper function to run a given process P with the Var process. If the parameter hide is true, runWith function hides all events introduced by the shared variable. In later chapters, we will see how the runWith function helps reduce the code complexity of the synchronization object implementation.

Figure 7 is an example of two processes using a shared variable. The first line in the example creates a shared variable VarA with value ranging from 0 to 2 and initialized with 0. Process P increments VarA modulo 3 forever and process Q reads VarA forever. Process P interleaves with process Q, and the combined process is further synchronized with the variable VarA process. In the resulting process System,

```
--set of possible value for the variable
--inital value for the variable
module ModuleVariable(TypeValue, initialValue)
  Var(value) = getValue!value → Var(value)
             □ setValue?value → Var(value)
  chanset = {|getValue, setValue|}
exports
  --(Bool, Proc) → Proc
  runWith(hide,P) = if hide then (Var(initialValue) [|chanset|] P) \ chanset
                    else  Var( initialValue ) [|chanset|] P
  channel getValue, setValue: TypeValue
endmodule

module ModuleUninitVariable(TypeValue)
  Var(value) = getValue!value → Var(value)
             □ setValue?value → Var(value)
  chanset = {|getValue, setValue|}
exports
  runWith(hide,P) =
    if hide then ((|~| x:TypeValue • Var(x)) [| chanset |] P) \ chanset
    else (|~| x:TypeValue • Var(x)) [| chanset |] P
  channel getValue, setValue: TypeValue
endmodule
```

Figure 6: The shared variable module in CSP

changes to VarA made by process P is visible to process Q.

```
instance VarA = ModuleVariable({0..2},0)
P = VarA::getValue?a → VarA::setValue!((a+1)%3) → P
Q = VarA::getValue?a → Q
System = VarA::runWith(false,P|||Q)
```

Figure 7: Example of two processes using a shared variable

2.2 Semaphore

A Semaphore is a simple but powerful concurrent primitive. This thesis shall describe and use a simplified binary semaphore from [Lea06], which removes interrupts and timeout operations.

A binary semaphore can either be raised or lowered. A up function call raises the semaphore regardless of the semaphore state. If a process calls the down method when the semaphore is raised, the semaphore becomes lowered. However, if the semaphore is unraised, the process waits until another process calls up and proceeds to put down the semaphore. Depending on the initial state of the semaphore, a binary semaphore can be further categorized as a mutex semaphore or a signaling semaphore.

Modeling a semaphore is straightforward in CSP. Figure 8 is the CSP semaphore module. A process may call up function or down function via channel upChan or channel downChan respectively. The semaphore is modelled by a process implemented by two mutually recursive functions Semaphore(True) and Semaphore(False). The semaphore process representing an unraised state accepts a upChan event by any process and proceeds to the raised process. The semaphore process representing a raised state can either accept a upChan event and recurse to the raised process, or accept a downChan event and proceed to the unraised process.

Like the shared variable in the earlier subsection, the semaphore is encapsulated in a CSP module. To create a semaphore, one needs to supply two arguments. TypeThreadID is the set of identities of processes that use this semaphore. initialState is a boolean value indicating the starting state of the semaphore. If initialState is true, the semaphore is raised initially. Otherwise, the semaphore is lowered.

```
  module ModuleSemaphore(TypeThreadID, initialState)
  --Raised
  Semaphore(cur) = if(cur) then downChan?id → Semaphore(False)
                           else upChan?id   → Semaphore(True)
  chanset = {|upChan, downChan|}
exports
  runWith(hide,P) = (Semaphore(initialState) [| chanset |] P)  \
                    (if hide then chanset else {})
  channel upChan, downChan: TypeThreadID
endmodule
```

Figure 8: The binary semaphore module in CSP

2.3 Monitor

2.3.1 JVM Monitor

A Monitor is another powerful concurrent primitive. This thesis will also use a simplified monitor from [Lea06]. JVM Monitor provides two key features, mutual exclusion and waiting.

Monitors can be used to prevent race conditions. At any time, only one process can run code inside a synchronized block that belongs to one monitor. The function

op1 in Figure 9 uses synchronized block to prevent race condition on variable a.

Inside a synchronized block, the process can also perform wait, notify, and notifyAll. When a process inside the synchronized block calls wait, the process suspends and waits for notification from other processes. Since a waiting process may be spurious waked up, so a wait call is used with a while loop and a condition. In the MoniorExample of Figure 9, op2 waits until there is 10 op1 calls. In op1, a process calls notifyAll after incrementing the shared variale a. When there aren't 10 op1 calls, process waiting in op2 first wakes, finds the condition $a < 10$ true, and returns to sleep.

```
class MonitorExample {
  private var a = 0;

  def op1():Unit = synchronized{
    a=(a+1)%20;
    notifyAll()
  }
  def op2():Unit = synchronized{
    while(a<10) wait();
  }
}
```

Figure 9: A simple Scala class that uses a monitor internally

2.3.2 Monitor Module

The monitor process has two states and behaves differently in two states, captured by two processes in Figure 11.

When there is no running process, the behavior of the monitor is captured by the CSP process inactive, with parameter waiting being the set of waiting processes. The monitor can allow a process to run synchronized code with a waitEnter event. Or, the monitor can spuriously wake a process with a SpuriousWake event, and the spuriously waken process behaves like a normally waken process.

When there is a running process, the behavior of the monitor is captured by the CSP process active, with parameter cur being the identity of the process running synchronized block, and waiting being the set of the waiting process. The monitor

process should respond to method calls from the running process cur, and the monitor should not allow another process to obtain the monitor lock. In active state, the monitor can also spuriously wake a waiting process.

On the client process side, most functions are implemented by simply synchronizing with the monitor on an event. For example, before entering the synchronized block, the process sends a WaitEnter event. The only exception is the wait function, as after being notified, the process needs to resume execution. The process first sends a wait event to tell the monitor that it is waiting and release the monitor lock. The monitor then receives a waitNotify or spuiousWake event, for being notified. Then the process reobtains the monitor lock with a waitEnter event.

The monitor module also provides a few useful macros. The synchronized function wraps a CSP code to run under the protection of the monitor. The whilewait function implements the common Scala pattern 'while(cond) wait()'. The implementation uses a functional replacement for while statement in imperative programming languages. The cond parameter is a CSP function of type '(Proc, Proc) -> Proc'. The return process of cond first performs some events to check the condition of the while statement. If the condition is true, the return process continues to run the process in the first parameter, or the return process runs the process in the second parameter.

With these process side functions and macros, the Scala MonitorExample class in Figure 9 can be converted into the CSP code in Figure 10.

There are two design choices worth mentioning in the implementation of the monitor.

First, note that both WaitNotify and SpuriousWake events come from the monitor process instead of directly synchronizing with the currently running process. When a process calls notify or notifyAll, it needs to synchronize with the monitor process. This is because the running process does not know how many processes are waiting. If the monitor is implemented the other way, the notifying process will block if there

```
instance VarA = ModuleVariable({0..20},0)
instance Monitor = ModuleMonitor(TypeThreadID, False)

op1(me)=synchronized(me,
  --a=(a+1)%20;
  VarA::getValue?x → VarA::setValue!((x+1)%20) →
  --notifyAll();
  Monitor::notifyAll(me)
)

op2(me)=synchronized(me,
  --while(a<10) wait()
  Monitor::whileWait(me, \ ktrue,kfalse •
    VarA::getValue?x → if x<10 then ktrue else kfalse
  )
)
```

Figure 10: The CSP implementation of MonitorExample in Figure 9

is no waiting process. Similarly, a process calling notifyAll does not know how many processes it should wake up.

Secondly, a monitor can introduce divergence by repetitively spuriously waking up a waiting process, whose condition keeps unsatisfied and never changes. This is an unwanted behavior in failure testing. So the monitor module has an extra parameter disableSpurious to disable spurious wakeups.

```
module ModuleMonitor(TypeThreadID)
  channel Notify, NotifyAll, Exit, Wait,
          WaitNotify, WaitEnter, SpuriousWake: TypeThreadID

  chanset = {| Notify, NotifyAll, Exit, Wait, WaitNotify, WaitEnter, SpuriousWake|}

  --A list of event for every event e in s
  repeat(ch, s) = if s={} then SKIP else ch?a:s → repeat(ch, diff(s, {a}))

  --cur is current active running thread
  --waiting is a set of threads waiting to be notified
  active(cur, waiting) =
    --current running thread notify
    Notify.cur → (
      --do nothing if no thread is waiting
      if waiting={} then active(cur, {})
      --wakeup a process
      else WaitNotify?a:waiting →
           active(cur, diff(waiting, {a}))
    ) □ --current running thread notifyAll
    NotifyAll.cur → (
      repeat(WaitNotify, waiting);
      active(cur, {})
    ) □ --current running thread exit
    Exit.cur → (
      inactive(waiting)
    ) □ --current running thread wait
    Wait.cur → (
      inactive(union(waiting,{cur}))
    ) □ --spurious wakeup
    waiting≠{} & SpuriousWake?a:waiting → (
      active(cur, diff(waiting, {a}))
    )

  --when no active thread is running
  inactive(waiting) =
    --pick a thread that is ready to enter
    WaitEnter?a → (
      active(a, waiting)
    ) □
    --spurious wakeup
    waiting≠{} & SpuriousWake?a:waiting → (
      inactive( diff(waiting, {a}))
    )
```

Figure 11: The CSP Monitor Module - Part 1 - the monitor process

```
exports
  --Given a process that uses the monitor
  --Return the process synchronized with the monitor server process
  --If hide is true, monitor channels are hidden
  runWith(hideSpurious, hideInternal, P) =
    let hideset0 = if hideInternal then chanset else {} within
    let hideset1 = if hideSpurious then hideset0
                   else diff(hideset0,{|SpuriousWake|}) within
    (inactive({}) [|chanset|] P) \ hideset1

  --java-like synchronized function
  synchronized(me, P)= WaitEnter.me → P; Exit.me → SKIP

  enter(me) = WaitEnter.me → SKIP

  exit(me) = Exit.me → SKIP

  --notify()
  notify(me) = Notify.me → SKIP

  --notifyAll()
  notifyAll(me) = NotifyAll.me → SKIP

  --wait()
  wait(me) =
    Wait.me → (
      (WaitNotify.me → WaitEnter.me → SKIP)
    □(SpuriousWake.me → WaitEnter.me → SKIP)
    )

  whileWait(me,cond) = while(cond)(wait(me);SKIP)
endmodule
```

Figure 12: The CSP Monitor Module - Part 2 - client process side functions

3 MenWomen

The MenWomen object is a classical problem from the concurrent programming course. In the problem, some processes need to pair off and share identities between the paired processes. Figure 13 is the interface of the MenWomen object. A process can call manSync from the object to pair with a process calling womanSync. Also, a process can call womanSync from the object to pair with another process calling manSync. For simplicity, if a process is calling manSync, we shall call it a man process. Similarly, a process calling manSync is called a woman process.

```
trait MenWomenT{
  def manSync(me: Int): Int
  def womanSync(me: Int): Int
}
```

Figure 13: Scala Interface of the MenWomen object

3.1 Implementation

One way to implement the MenWomen object is to use a monitor and a shared variable indicating the stage of synchronization. Figure 14 is a Scala implementation of the MenWomen object with monitor.

- A man process enters the synchronization and waits until the current stage is 0. Then in stage 0, the man process sets the global variable him inside the MenWomen object to its identity. Then the man process notifies all processes so that a waiting woman process can continue. Finally, the man process waits for stage 2.

- A women process enters the synchronization and waits until the current stage is 1. The woman process sets the global variable her to its identity and returns the value of the global variable him.

- In stage 2, the waiting man process in stage 0 is wakened up by the woman

process in stage 1. The man process notifies all waiting processes and returns the value of her.

The code snippet in Figure 14 is a Scala implementation of the MenWomen process using a monitor. With the shared variable and monitor module, the Scala code is further translated to a CSP code in Figure 15. With the convention described in the introduction section, every function call begins with a Call event containing all parameters. And every function call ends with a Return event containing the return value.

```
class MenWomen extends MenWomenT{
  private var stage = 0
  private var him = −1
  private var her = −1

  def manSync(me: Int): Int = synchronized{
    while(stage != 0) wait()
    him = me; stage = 1; notifyAll()
    while(stage != 2) wait()
    stage = 0; notifyAll(); her
  }

  def womanSync(me: Int): Int = synchronized{
    while(stage != 1) wait()
    her = me; stage = 2; notifyAll();
  }
}
```

Figure 14: A correct MenWomen object implementation in Scala

3.2 Linearization Test

Recall that in the testing system, each process can call any function provided by the concurrent datatype or choose to terminate. In defining processes in the testing system, a helper function is used. chaosP(P) runs the process P forever or terminates after running a finite number of P. The processes in the testing system simply use chaosP with a process that non deterministically chooses to perform manSync or womanSync with their identities.

```
chaosP(P) = (P;chaosP(P)) ⊓ SKIP
Thread(me)=chaosP(manSync(me) ⊓ womanSync(me))
```

```
instance VarStage = ModuleVariable({0,1,2},0)
instance VarHim = ModuleUninitVariable(TypeThreadID)
instance VarHer = ModuleUninitVariable(TypeThreadID)
instance Monitor = ModuleMonitor(TypeThreadID)
manSync(me) =
  Call ! me ! ManSync →
  Monitor::enter(me);
    Monitor::whileWait(me, \ ktrue,kfalse •
      VarStage::getValue ? x →
      if x≠0 then ktrue else kfalse
    );
    VarHim::setValue ! me →
    VarStage::setValue ! 1 →
    Monitor:: notifyAll (me);
    Monitor::whileWait(me, \ ktrue,kfalse •
      VarStage::getValue ? x →
      if x≠2 then ktrue else kfalse
    );
    VarStage::setValue ! 0 →
    Monitor:: notifyAll (me);
    VarHer::getValue ? ans →(
  Monitor:: exit (me);
  Return ! me ! ManSync ! ans→
  SKIP
  )
womanSync(me)=
  Call ! me ! WomanSync →
  Monitor::enter(me);
    Monitor::whileWait(me, \ ktrue,kfalse •
      VarStage::getValue ? x →
      if x≠1 then ktrue else kfalse
    );
    VarHer::setValue ! me →
    VarStage::setValue ! 2 →
    Monitor:: notifyAll (me);
    VarHim::getValue ? ans →(
  Monitor:: exit (me);
  Return ! me ! WomanSync ! ans→
  SKIP
  )
```

Figure 15: Translated CSP code for the correct MenWomen object immplementation

All processes in the testing system interleave with other processes and synchronize with the processes of shared variables and the process of the monitor. Since the testing system should only include Call and Return events, all other events are hidden using the first two boolean flags in runWith defined in earlier sections.

```
System(All)=runWith(True,True,||| me:All • Thread(me))
```

Sync events represent the synchronization between a man process and a woman process, which is used internally in the testing specification. The CSP code below is the definition and specification of Sync channel.

```
channel Sync: TypeThreadID . TypeThreadID . TypeThreadID . TypeThreadID
Spec(All) = Sync ? man:All ? woman:diff(All,{man}) ! woman ! man → Spec(All)
```

A Sync event takes four parameters, the identity of the man process, the return value of the man process, the identity of the woman process, and the return value of the woman process. For each synchronization, the spec process ensures two properties. First, synchronization occurs between two different processes. Second, the return value of each participating process is the identity of the other participating process.

A linearizer ensures two properties. The return value of a function call comes from the synchronization. And the synchronization point occurs sometime during the function call. The Linearizer function for MenWomen achieves this by two branches of CSP processes, each representing a method in the object. Each branch starts with a Call event, then a Sync event with its identity, and finally a Return event which computes from the Sync event. In addition to the two branches, the linearizer process should be able to STOP to match the behavior in the testing system.

Finally, Linearizer(All) composites the Sync specification with the linearizer processes so that the Sync events are valid according to the Sync specification and Sync events are agreed by both participating process. The latter property is achieved using replicated generalized parallel, described in [GRPAR13]. By the semantic of the replicated generalized parallel, when a process wants to perform an event, the process must synchronize all other processes which can perform this event. Figure

16 visualizes a trace generated by 1,T2Linearizers(T) by the FDR debugging window. When a linearizer performs Call and Return, the linearizer does not synchronize with any process. When linearizer(T1) performs Sync.T1.T2.T2.T1, linearizer(T2) and Spec must also perform the Sync event.

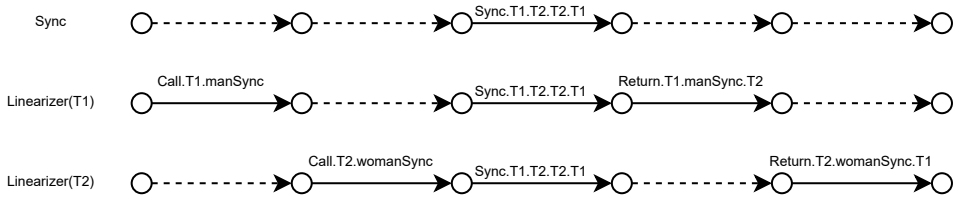

Figure 16: FDR Visualization of traces of a man process and a woman process synchronizing

All above properties suffice to show that Linearizers(All) generates all traces that correspond to a valid history and that Linearizers is a valid specification process for the testing system. Informally, for any valid history, a linearizability test can find a set of synchronization points, which corresponds to a trace of the specification for Sync events. Furthermore, the valid history can be obtained by adding respective call and return events around synchronization points.

```
Lin(All,me)= (
  Call ! me!ManSync→
  Sync!me?mereturn?other:diff(All,{me})?otherreturn →
  Return!me!ManSync!mereturn →
  Lin(All,me)
)⊓(
  Call ! me!WomanSync →
  Sync?other:diff(All,{me})?otherreturn!me?mereturn →
  Return!me!WomanSync!mereturn →
  Lin(All,me)
)⊓STOP
LinEvents(All,me)=union({
  ev | ev←{|Sync|},
  let Sync.t1.a.t2.b=ev within
    countList(me,<t1,t2>)=1 and member(t1,All) and member(t2,All)
},{|Call.me,Return.me|})

Linearizers(All)=((|| me: All • [LinEvents(All,me)] Lin(All,me)) [|{|Sync|}|] Spec(All))
                  \ {|Sync|}
```

Figure 17: Definition of linearizer process in CSP

Finally, we perform the test using trace refinement for safety property and failure

refinement for liveness. As expected, the correct implementation passes all tests.

```
System2=System({T1,T2})
Spec2Thread=Linearizers({T1,T2})
assert Spec2Thread ⊑_T System2
assert Spec2Thread ⊑_F System2
```

3.3 A faulty version

We shall examine another MenWomen implementation in Figure 18. One key difference in this faulty MenWomen object is that it uses Option data in the shared variables to store the identity of the process calling manSync and the process calling womanSync.

```
class FaultyMenWomen extends MenWomenT{
  private var him: Option[Int] = None
  private var her: Option[Int] = None

  def manSync(me: Int): Int = synchronized{
    while(him.nonEmpty) wait()
    him = Some(me); notifyAll()
    while(her.isEmpty) wait()
    val Some(res) = her
    her = None; notifyAll()
    res
  }

  def womanSync(me: Int): Int = synchronized{
    while(her.nonEmpty) wait()
    her = Some(me); notifyAll()
    while(him.isEmpty) wait()
    val Some(res) = him
    him = None; notifyAll()
    res
  }
}
```

Figure 18: A Faulty Scala implementation of MenWomen object

The safeness property test shows that this implementation handles scheduling carelessly. This implementation fails the test Spec2Thread $\sqsubseteq_T$ System2, and FDR provides a trace that violates the safeness specification, shown in Figure 19. FDR further allows the user to expand the hidden τ events in the testing system, which are normally processes' interactions with shared variables and the monitor. By understanding the expanded trace in CSP, we can find an equivalent way to trigger the bug in Scala.

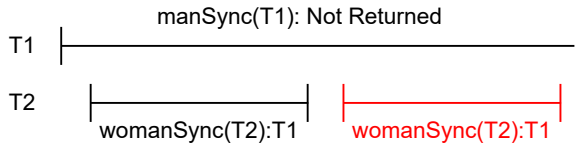

Figure 19: Trace

- A process T1 calls manSync. On first line, since the shared variable him is initially None, T1 skips the wait, set him to Some(T1) and waits for a process calling womanSync.
- A process T2 calls womanSync and returns T1. At this stage, there is no waiting process waiting to run womanSync and the shared variable is not None. So T2 does not wait at any point, notifies all waiting process, and returns.
- Before T1 reenters the synchronized block, process T2 calls womanSync again. him has not been reset by T1 yet. So T2 pairs with T1 again, which should not be allowed.

4 ABC

With an ABC object, three processes can exchange data with each other two processes. More specifically, one process calling aSync, one process calling bSync, and one process calling cSync synchronizes. Then each of the three processes returns with the arguments of two other processes. For simplicity, we shall call a process calling syncA as an A-process, a process calling syncB as a B-process, and a process calling syncC as a C-process.

One of the challenges to check an ABC object is the huge number of states in the CSP model. In this section, we shall see how the linearizer process can be optimized and the efficiency of the explicit linearization point test.

4.1 Implementation

Figure 20 is a Scala implementation of a ABC object with semaphore. In each round of synchronization,

- Initially semaphore aClear is raised. An A-process acquires semaphore aClear, sets the shared variable a to its parameter, raises semaphore bClear and waits to acquire semaphore aSignal. A B-process and a C-process behaves similarly in turn, except they use different semaphores and shared variables.

- After a C-process raises semaphore aSignal, the A-process is able to continue. The A-process reads the shared variable b and c, raises the semaphore bSignal, and returns b and c. Likewise, B and C also take the value of two other shared variable and raise respective semaphores in turn.

Using the shared variable and semaphore module, it is easy to translate the Scala implementation to a CSP implementation.

```
class ABC[A,B,C] extends ABCT[A,B,C]{
  // The identities of the current (or previous) threads.
  private var a: A = _
  private var b: B = _
  private var c: C = _

  // Semaphores to signal that threads can write their identities.
  private val aClear = MutexSemaphore()
  private val bClear, cClear = SignallingSemaphore()

  // Semaphores to signal that threads can collect their results.
  private val aSignal, bSignal, cSignal = SignallingSemaphore()

  def syncA(me: A) = {
    aClear.down         // (A1)
    a = me; bClear.up // signal to b at (B1)
    aSignal.down        // (A2)
    val result = (b,c)
    bSignal.up          // signal to b at (B2)
    result
  }

  def syncB(me: B) = {
    bClear.down         // (B1)
    b = me; cClear.up // signal to C at (C1)
    bSignal.down        // (B2)
    val result = (a,c)
    cSignal.up          // signal to c at (C2)
    result
  }

  def syncC(me: C) = {
    cClear.down         // (C1)
    c = me; aSignal.up // signal to A at (A2)
    cSignal.down        // (C2)
    val result = (a,b)
    aClear.up           // signal to an A on the next round at (A1)
    result
  }
}
```

Figure 20: A semaphore-based Scala implementation of the ABC object

```
instance VarA = ModuleUninitVariable(TypeData)
instance VarB = ModuleUninitVariable(TypeData)
instance VarC = ModuleUninitVariable(TypeData)

instance aClear = ModuleMutexSemaphore(TypeThreadID)
instance bClear = ModuleSignallingSemaphore(TypeThreadID)
instance cClear = ModuleSignallingSemaphore(TypeThreadID)
instance aSignal = ModuleSignallingSemaphore(TypeThreadID)
instance bSignal = ModuleSignallingSemaphore(TypeThreadID)
instance cSignal = ModuleSignallingSemaphore(TypeThreadID)

runWith(hide,p)=
  VarA::runWith(hide,
  VarB::runWith(hide,
  VarC::runWith(hide,
  aClear::runWith(hide,
  bClear::runWith(hide,
  cClear::runWith(hide,
  aSignal::runWith(hide,
  bSignal::runWith(hide,
  cSignal::runWith(hide,
    p
  )))))))))

SyncA(me,avalue) =
  Call ! me ! ASync ! avalue →
  --aClear.down
  aClear::downChan ! me →
  --a = me
  VarA::setValue ! avalue →
  --bClear.up
  bClear::upChan ! me →
  --aSignal.down
  aSignal::downChan ! me →
  --(b,c)
  VarB::getValue ? b →
  VarC::getValue ? c →
  --bSignal.up
  bSignal::upChan ! me →
  --result →
  Return ! me ! ASync ! (b.c) →
  SKIP

...
```

Figure 21: Translated CSP Code for the correct ABC implementation

4.2 Testing

Similar to the MenWomen object, a testing system is defined through any number of working processes and a specification is built from a Sync channel, linearizer processes, and a synchronization alphabet set.

One key difference of the ABC object is that processes can call with any argument from the set TypeData, whereas in MenWomen object, processes can only call with their identity. Inside chaosP, the processes also choose an argument with General Non-Deterministic Choice.

```
Thread(me)=chaosP( ⊓ x:TypeData • (
    SyncA(me,x)
  ⊓ SyncB(me,x)
  ⊓ SyncC(me,x)
))
```

The testing specification uses the same component, the Sync channel, linearizer processes, and linEventns. The definition of Sync channel is shown in Figure 22. The event $Sync.t_1.a.b.c.t_2.d.e.f.t_3.g.h.i$ represents the synchronizations between three threads, t_1, t_2, t_3. Process t_1 calls aSync with a and returns (b, c). The second process t_2 calls bSync with d and returns (e, f). And the last process t_3 calls cSync with g and returns (h, i). The Sync specification process, shown in the same figure, checks that in each Sync events, the return value of each process is the pair of arguments of the two other function call.

Figure 23 is the definition of a linearizer process, written a similar format to the MenWomen object.

```
--thread identity calling ASync. ASync parameter a. ASync return pair (b,c)
--thread identity calling BSync. BSync parameter b. BSync return pair (a,c)
--thread identity calling CSync. CSync parameter c. CSync return pair (a,b)
channel Sync: TypeThreadID.TypeData.TypeData.TypeData.
              TypeThreadID.TypeData.TypeData.TypeData.
              TypeThreadID.TypeData.TypeData.TypeData

Spec(All) =
  Sync?t1:All?a?b?c
      ?t2:diff(All,{t1})!b!a!c
      ?t3:diff(All,{t1,t2})!c!a!b →
  Spec(All)
```

Figure 22: Definition of Sync channel and specification of Sync event

```
--Linearizer for a process
Lin(All,me)=(
  --me synchronizes as thread A
  Call ! me!ASync?a →
  Sync!me!a?b?c?t2:diff(All,{me})?t2b?t2a?t2c?t3:diff(All,{me,t2})?t3c?t3a?t3b →
  Return!me!ASync!b!c →
  Lin(All,me)
) ⊓ (
  --me synchronizes as thread B
  Call ! me!BSync?b →
  Sync?t2:diff(All,{me})?t2b?t2a?t2c!me!b?a?c?t3:diff(All,{me,t2})?t3c?t3a?t3b →
  Return!me!BSync!a!c →
  Lin(All,me)
) ⊓ (
  --me synchronizes as thread C
  Call ! me!CSync?c →
  Sync?t2:diff(All,{me})?t2b?t2a?t2c?t3:diff(All,{me,t2})?t3a?t3b!me!c?a?b →
  Return!me!CSync!a!b →
  Lin(All,me)
)
```

Figure 23: Definition of linearizer process

```
System3=System({T1,T2,T3})
Spec3Thread=Linearizers({T1,T2,T3})

assert Spec3Thread ⊑_T System3
assert Spec3Thread ⊑_F System3
```

4.2.1 Speeding up model compilation

Consider the specification process with three processes. Let M be the size of the set of all possible arguments. Conside r the trace in Figure 24, where process T1 calls aSync with A, T2 calls bSync with B, T3 calls cSync with C. Then they synchronization.

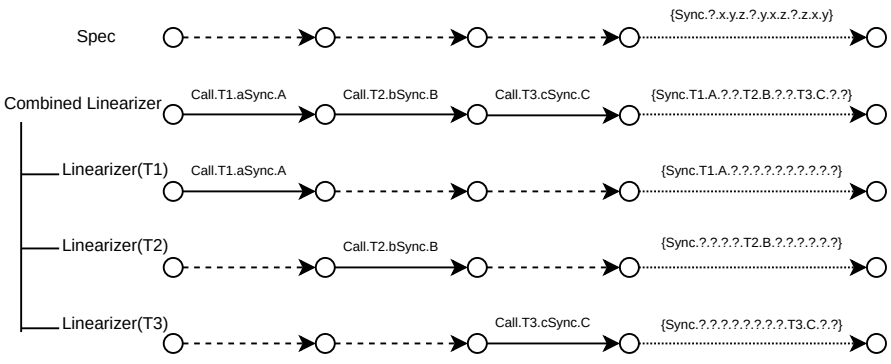

Figure 24: The set of possible Sync event for each CSP process

The last transition in the diagram is the Sync event between three processes.

Above the edge is the set of possible Sync event that every process accepts. Each linearizer accepts $3^2 * M^8$ possible Sync event. The combined linearizer accepts M^6 sync event. However, according to the sync specification, only one Sync event is valid.

With the above analysis, it is tempting to reduce the redundancy in Sync event. Optimize the linearizer process by using the information from the specification process. Instead of choosing all possible remaining arguments, the individual linearizer could choose correct arguments according to the specification process. Figure 25 includes part of the simplified code. This change does not reduce the number of transitions in the resulting specification, but it helps FDR build the process faster.

With this optimization, the testing for less than 5 processes finishes quickly.

```
--A modified version of lin2
--slow
channel Sync: TypeThreadID . TypeThreadID . TypeThreadID . TypeData . TypeData . TypeData

--Linearizer for a process
Lin(All,me)= (
      --me synchronizes as thread A
      Call ! me!ASync?a →
      Sync!me?t2:diff(All,{me})?t3:diff(All,{me,t2})!a?b?c →
      Return!me!ASync!b!c →
      Lin(All,me)
   ) ⊓ (
      --me synchronizes as thread B
      Call ! me!BSync?b →
      Sync?t2:diff(All,{me})!me?t3:diff(All,{me,t2})?a!b?c →
      Return!me!BSync!a!c →
      Lin(All,me)
   ) ⊓ (
      --me synchronizes as thread C
      Call ! me!CSync?c →
      Sync?t2:diff(All,{me})?t3:diff(All,{me,t2}) !me?a?b!c →
      Return!me!CSync!a!b →
      Lin(All,me)
   ) ⊓ STOP
LinEvents(All,me)=union({
  e | e ← {|Sync|},
  let Sync.t1.t2.t3.a.b.c=e within
    countList(me,<t1,t2,t3>)=1 and member(t1,All) and member(t2,All) and member(t3,All)
},{|Call.me,Return.me|})
Linearizers(All)= (|| me: All • [LinEvents(All,me)] Lin(All,me)) \ {|Sync|}
```

Figure 25: Simplified definition of Sync channel and part of simplified linearizer

4.3 Faulty version

Recall that in Java and Scala, raising a semaphore immediately allows another thread waiting to acquire the semaphore to continue. So it is essential to take a copy of the two other arguments before raising the semaphore.

On the other hand, what if the implementation of syncA does not take a copy of the argument? It turns out that the faulty ABC object passes tests for three processes but fails the linearisation test with at least four threads.

4.3.1 Explanation of the error case

For the test Spec4Thread $\sqsubseteq_T$ System4, FDR displays a trace of the testing system that violates the specification. From the trace, it seems that process T1 synchronizes with T2 and T3 in the first round, and should return (B,C), but (E,F), the argument in the second round is returned. Expanding the τ event and translating CSP traces into program traces makes it possible to see what goes wrong in the faulty version when there are four threads.

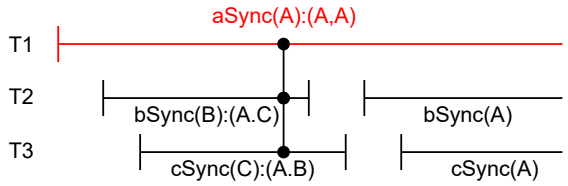

Figure 26: Trace that violates the specification

- In the first round of synchronization, process T_A, T_B, T_C call aSync, bSync or cSync respectively, and put down its argument in turn.
- Process T_A raises bSignal. Before T_A exits, the other two processes T_B, T_C returns. Now T_A should return argument of T_B and T_C.
- Another round of synchronization starts. Thread T_D, T_B, T_C call aSync, bSync or cSync respectively, and overwites the shared variable a,b,c in turn.

- Now T_A returns with (b, c) from the second round, which may not be the argument of bSync and cSync in the first round.

4.4 Explicit signaling point test

In this section, we describe a faster testing technique. In the explicit signal test, one first observes the objects and makes a hypothesis on the order of signal points, where one process sends data to another process. Then one uses CSP model to check signal points actually signal in the hypothsised order, and correspond to each of the calls and returns.

In the correct implementation of the ABC object, six signal points should occur in order during each round of synchronization. The A-process, the B-process, and the C-process set the shared variable a, b, and c in turn before raising respective semaphore. Then the A-process, B-process, and the C-process read the shared variable in turn after lowering the respective semaphore.

To check if the signal points actually occur in order, we perform a trace refinement test. In the implementation, we replace every piece of signaling code with a signal event. For example, in SyncA shown in Figure 27, a=me is replaced with a SP1 event, and (b,c) is replaced with a SP4 events. The variable modules are also modified to support reading and writing on corresponding signal point events.

The specification process, shown in Figure 28, is a process that repeats the six signal point events in order, with the extra constraint that the last three SP events match the first three SP events.

Immediately we can see the specification process is much simpler than the specification process in linearization testing. But how fast is the explicit synchronization test? We use the number of states and transitions reported by FDR as indicators of the complexity of the process and test. Specifically the test Spec $\sqsubseteq_T$ Spec and System $\sqsubseteq_T$ System indicate the complexity of the specification process and the testing system. And the test Spec $\sqsubseteq_T$ System indicates the complexity of the actual

```
SyncA(me,avalue) =
  −−aClear . down
  aClear :: downChan ! me →
  −−a = me
  SP1 ! me ! avalue →
  −−bClear . up
  bClear :: upChan ! me →
  −−aSignal . down
  aSignal :: downChan ! me →
  −−(b,c)
  SP4 ! me ! bvalue ! cvalue →
  −−bSignal . up
  bSignal :: upChan ! me →
  SKIP
```

Figure 27: aSync function in explicit linearization point testing

```
spec(All) =
  SP1 ? t1:All ? a →
  SP2 ? t2:diff(All,{t1}) ? b →
  SP3 ? t3:diff(All,{t1,t2}) ? c →
  SP4 ! t1 ! b ! c →
  SP5 ! t2 ! a ! c →
  SP6 ! t3 ! a ! b →
  spec(All)
```

Figure 28: The specification process of explicit signal point test

testing. We use three processes in the test and the statistics are organized in the table below.

	State		Transition	
	Linearization	Explicit signal	Linearization	Explicit signal
Spec $\sqsubseteq_T$ Spec	$1.02 * 10^5$	$1.75 * 10^2$	$2.91 * 10^5$	$2.23 * 10^2$
System $\sqsubseteq_T$ System	$3.81 * 10^5$	$1.24 * 10^5$	$1.65 * 10^6$	$5.60 * 10^5$
Spec $\sqsubseteq_T$ System	$3.81 * 10^5$	$1.24 * 10^5$	$1.65 * 10^6$	$5.60 * 10^5$

As we can see from the table, the state machine of the specification process is much smaller in the explicit synchronization test. It is no surprise because the specification process in the explicit synchronization test is a simple recursive process. The state machine of the testing system is slightly smaller because the Call and Return events are removed. In the linearization test, the Return events can be delayed and occur in the trace of later synchronization, which can make the testing system more complicated. The simpler specification and testing system makes the explicit signal test fast.

However, the drawback of the explicit synchronization test is that it needs to be done on a per implementation basis. In the faulty implementation, the signal events no longer occur in the same order. As a result, a new test needs to be written.

5 Terminating Queue

A terminating queue provides thread-safe enqueue and dequeue operations. Suppose a process dequeues when the queue is empty. The process blocks and waits for some process to enqueue. In addition, if all processes are dequeuing, no new element can be added to the queue and the whole system deadlocks. In this case, the queue terminates and returns None for every process. Figure 29 is the interface of a terminating queue.

```
trait TerminatingQueueT[A]{
  def enqueue(x: A): Unit
  def dequeue: Option[A]
}
```

Figure 29: Scala interface of terminating queue

A terminating queue is a stateful synchronization object, as earlier enqueue and dequeue operations affect later synchronizations. The timeline diagram below is invalid because the second dequeue call should return 1 instead of 2.

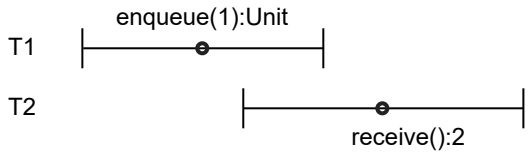

Figure 30: A faulty timeline of termianting queue

5.1 Implementation

The Scala implementation of the Terminating Queue wraps a Scala Queue with a monitor. Upon creation, the object is given numWorkers, the number of processes using the terminating queue. In addition to an internal queue queue, the terminating queue uses two more shared variables. The shared variable waiting is the number of processes waiting to dequeue an element from the queue. The shared variable done indicates if the queue has terminated. If the value of done is true, then any further function call should do nothing.

Enqueue is a trivial operation. In the synchronized block, the process adds the element to the internal queue and notifies a process waiting to dequeue an element. Dequeuing is also trivial when the internal queue is not empty. The process simply performs dequeue on the internal queue. When the queue is empty, if the value of waiting is numWorkers−1, then all processes are now waiting to dequeue, so the queue should terminate. The process notifies all waiting processes and returns None. Otherwise, the process increments the counter waiting and waits. The waiting process may be wakened for two cases. In the first case, a new element is added to the queue. The process decrements waiting and returns the queue head as normal. In the second case, the process is wakened because the queue is terminating, and the process should return None.

```
class TerminatingQueue[A](numWorkers: Int){
  private val queue = new Queue[A]
  private var waiting = 0
  private var done = false

  def enqueue(x: A) = synchronized{
    if(!done){
      queue.enqueue(x)
      if(waiting > 0) notify()
    }
  }

  def dequeue: Option[A] = synchronized{
    if(!done && queue.isEmpty){
      if(waiting == numWorkers−1){ // System should terminate
        done = true; notifyAll()
      }
      else{
        waiting += 1
        while(queue.isEmpty && !done) wait()
        waiting −= 1
      }
    }
    if(done) None else Some(queue.dequeue)
  }
}
```

Figure 31: Scala implementation of terminating queue

A few workarounds are required to keep the state machine of the CSP model finite and analyzable for FDR. First, the range of variable waiting is limited to integers from 0 to numWorkers inclusive, instead of all 2^{32} possible values from a Scala integer. However, while building the state machine, FDR attempts to set the value of waiting to −1 and numWorkers+1 and throws an error. Even though the value of waiting

will never be −1 or numWorkers+1 in the correct implementation, FDR is unable to derive such information while building the state machine for dequeue when waiting is 0 or numWorker. As a workaround, we guard waiting settings with an if statement in Figure 33. Before setting the value of waiting, the process checks if the new value is in range. If not, the process diverges. In the testing system, we check that system is divergence free to show that waiting is always in range as expected.

Secondly, the Scala implementation uses a Scala Queue internally, which has unlimited capacity and cannot be modelled with finite states in CSP. To address this, we use a capacity-limited queue which behaves the same as the infinite queue when the queue is not full. However, if a process enqueues a new element when the internal queue is full, the process STOP or DIV according to what is required for testing. In Figure 32, we implement the capacity-limited queue with many functions, The interface may be slightly overkilling for a capacity-limited queue, but it is necessary for a more powerful finite-state queue described in later subsections.

```
--Maximum number of element in the queue
NQueueCapacity=3
--Set of states of the capacity-limited queue
TypeQueue={q|i<-{0..NQueueCapacity},q<-ArrangementInList(i,TypeData)}
--Returns if the queue is empty
qEmpty(q)=q=<>
--Returns an empty queue
qNewQueue=<>
--Given a queue q, return a set of element that can be enqueued to queue q
qValidEnqueue(q)=if length(q)≥NQueueCapacity then {} else TypeData
--Given an element x, return a set of queues that accept an enqueue of x
qCanEnqueue(x)={q|q<-TypeQueue, length(q)<NQueueCapacity}
--Return the state of the queue after enqueueing x
qEnqueue(q,x)=q^<x>
--Return the set of all possible result and state pair after dequeueing
qDequeue(<>)={}
qDequeue(<qhead>^qtail)={(qhead,qtail)}
```

Figure 32: CSP Implementation of capacity-limited queue

5.2 Linearization Testing

We use similar approaches to construct the testing system and the specification. For the terminating queue object, there are three synchronizations between a process and the object. When a process enqueues or successfully dequeues, it synchronizes with

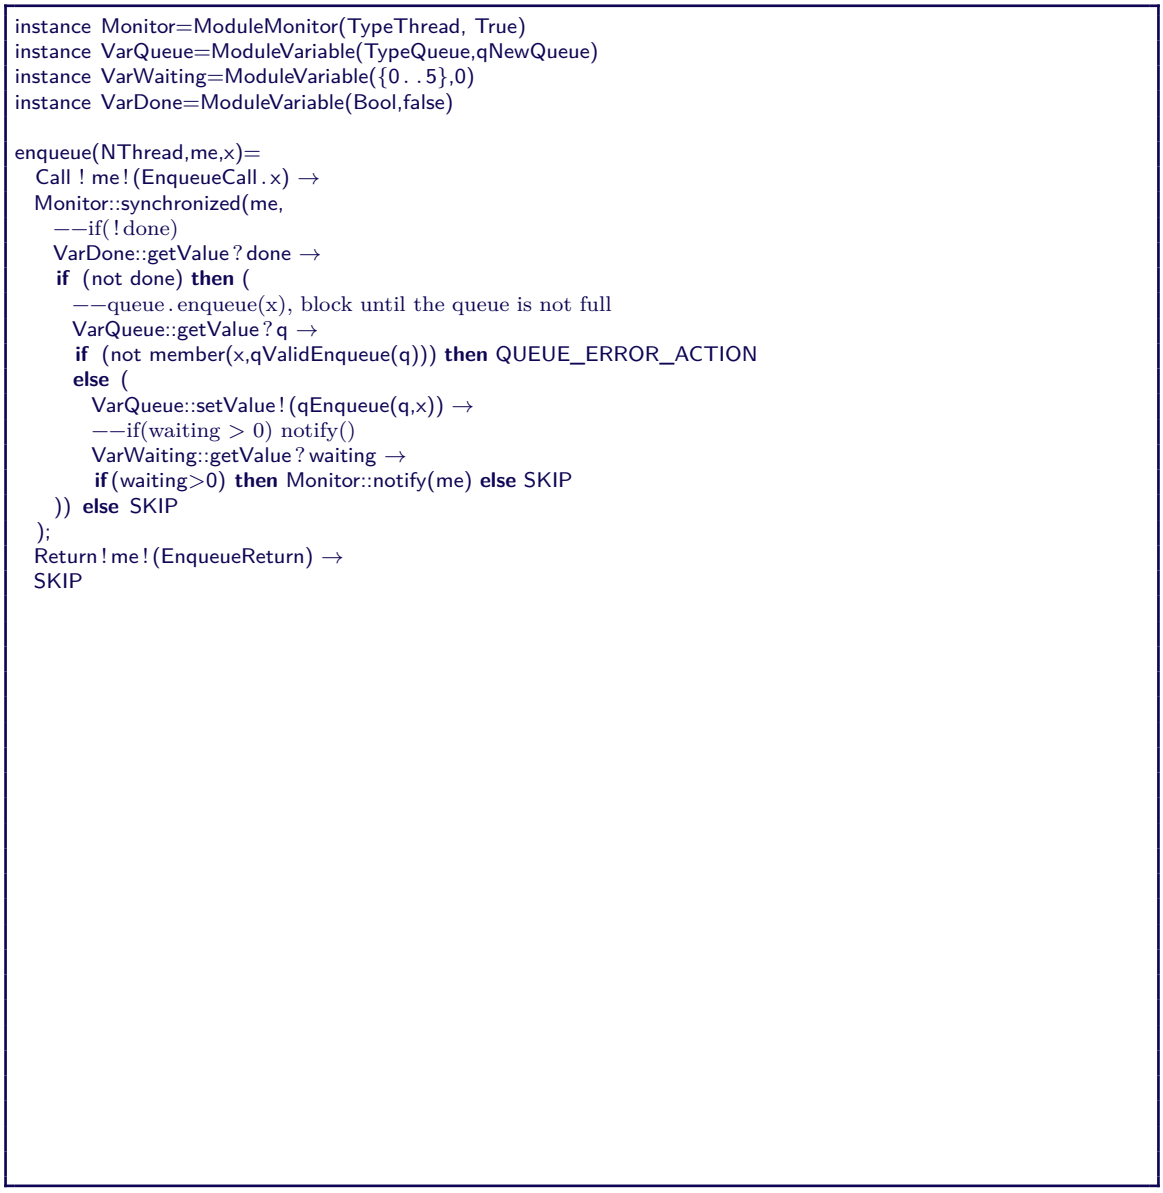

```
instance Monitor=ModuleMonitor(TypeThread, True)
instance VarQueue=ModuleVariable(TypeQueue,qNewQueue)
instance VarWaiting=ModuleVariable({0..5},0)
instance VarDone=ModuleVariable(Bool,false)

enqueue(NThread,me,x)=
  Call ! me ! (EnqueueCall . x) →
  Monitor::synchronized(me,
    --if(! done)
    VarDone::getValue ? done →
    if (not done) then (
      --queue . enqueue(x), block until the queue is not full
      VarQueue::getValue ? q →
      if (not member(x,qValidEnqueue(q))) then QUEUE_ERROR_ACTION
      else (
        VarQueue::setValue ! (qEnqueue(q,x)) →
        --if(waiting > 0) notify()
        VarWaiting::getValue ? waiting →
        if (waiting>0) then Monitor::notify(me) else SKIP
    )) else SKIP
  );
  Return ! me ! (EnqueueReturn) →
  SKIP
```

Figure 33: CSP implementation of the termianting queue - Enqueue Part

```
dequeue(NThread,me)=
  Call ! me ! (DequeueCall) →
  Monitor::enter(me);
  ––if( ! done && queue . isEmpty)
  VarDone::getValue ? done →
  VarQueue::getValue ? q →
  ( if (not done and qEmpty(q)) then (
    ––if(waiting = numWorkers−1)
    VarWaiting::getValue ? x→
    if (x≥NThread) then DIV
    else if (x=NThread−1) then(
      ––done=true
      VarDone::setValue ! true→
      ––notifyAll
      Monitor:: notifyAll (me)
    )else(
      ––waiting+=1
      ––if(x=NThread) then div else
      VarWaiting::setValue ! (x+1) →
      ––while(queue . isEmpty && ! done) wait()
      Monitor::whileWait(me, \ ktrue,kfalse •
        VarQueue::getValue ? q →
        VarDone::getValue ? done →
        if (qEmpty(q) and not done) then ktrue else kfalse
      );
      ––waiting −= 1
      VarWaiting::getValue ? y→
      if y=0 then DIV else
      VarWaiting::setValue ! (y−1)→
      SKIP
    )
  )else SKIP);
  ––if(done) None else Some(queue . dequeue)
  VarDone::getValue ? done →
  if (done) then (
    Monitor::exit (me);
    Return ! me ! (DequeueReturnNone) →
    ––the thread should stop any work
    STOP
  ) else (
    VarQueue::getValue ? q →
    if qEmpty(q) then DIV else
    □ (ans, qtail ): qDequeue(q) •
    VarQueue::setValue ! qtail →
    Monitor::exit (me);
    Return ! me ! (DequeueReturn . ans) →
    SKIP
)

runWith(hideSpurious,hide,P)=
  VarQueue::runWith(hide,
  VarWaiting::runWith(hide,
  VarDone::runWith(hide,
  Monitor::runWith(hideSpurious,hide,
    P
  ))))
```

Figure 34: CSP implementation of the termianting queue - Dequeue Part

the object and acts on the internal queue. For enqueue and dequeue, the synchronization is represented by a Sync event with the identity of the process, an object representing the function call, and the return value. When the queue shuts down and a process returns None, the process synchronizes with all other processes. In this case, the synchronization is represented by SyncShutdown.

Since a terminating queue is a stateful synchronization object, the specification of Sync event is different from the earlier Sync specification, in that the specification uses a parameter to keep the current queue. For both queues, the Sync specification should ensure that elements are added and removed in a First-In-First-Out order, and all enqueue and dequeue operation are valid for the queue. Sync specification does not check SyncShutdown however. Figure 35 shows the definition of Sync event and the specification.

```
channel Sync: TypeThread . TypeCallParam . TypeReturnParam
channel SyncShutDown

Spec(q)=
  (qValidEnqueue(q)≠{} & □ x:qValidEnqueue(q) • (
    Sync ? t ! (EnqueueCall . x) ! EnqueueReturn →
    Spec(qEnqueue(q,x))
  ))□
  (qDequeue(q)≠{} & □(x,newq):qDequeue(q) • (
    Sync ? t ! (DequeueCall) ! (DequeueReturn . x) →
    Spec(newq)
  ))□(
    qEmpty(q) & SyncShutDown → STOP
  )
```

Figure 35: Definition of Sync event and the specification

In addition to the safeness and liveness we usually test, we first check that the shared variable waiting is actually in the range 0 to numWorkers. We set spurious wakeup visible and invalid enqueue to cause STOP. If the testing system never diverges, it must be that a process always set a value from 0 to numWorkers for the variable waiting.

For safeness and liveness testing, we use the normal setup, in which the spurious wakeup is invisible and the process diverges when it performs an invalid enqueue operation.

5.3 Faulty Implementation

In this section, we use the linearization test to distinguish the correct implementation with three faulty implementations. In the first faulty implementation, enqueue operation always adds A to the internal queue regardless of the parameter of enqueue. In the second faulty implementation, enqueue method notify does not wake up a process waiting to dequeue. The third implementation has a miss-by-one error. In dequeue, A process checks waiting = numWorker before incrementing waiting. The result is shown in the table below.

	Divergence Free	Trace Refinement	Failure Refinement
Correct	Pass	Pass	Pass
Faulty1	Pass	Fail	Fail
Faulty2	Pass	Pass	Fail
Faulty3	Pass	Pass	Fail

With no surprise, the correct implementation passes all tests. The first faulty implementation fails all tests. The second faulty implementation fails only when spurious wakeups are disabled, as the missing notify can be compensated by some "coincidental" spurious wakeup. It shows that when testing an object that internally uses a monitor, one should test it with spurious wake and without spurious wakeups from the monitor. The third faulty implementation shows a similar pattern. Because of the miss-by-one error, the last process will not terminate the queue. As a result, all processes refuse to return. However, as spurious wakeups can cause divergence, the failure model is unable to capture this error.

5.4 Test with another queue

In this section, we perform a more general test on the correct and faulty implementation by using an unbounded queue. We used the A*BC* queue trick from [Low22]. The A*BC* queue does not have a capacity limit, but only accepts enqueuing of

some A, a single B, and some C.

The trick is applicable because the terminating queue is data independent. The terminating queue only stores elements in the internal queue and does not perform any operation on the element received.

There are three properties we can check from A*BC* terminating queue. First, the terminating queue should not duplicate any element. Otherwise, B may be duplicated, which fails refinement testing. Secondly, the terminating queue should not miss any element. Otherwise, B may be missed. Finally, when the terminating queue holds B, the queue must be non-empty. Thus the queue should not terminate in this case.

To implement the A*BC* terminating queue, all we need to do is to implement a regular A*BC* queue and import the regular A*BC* in the A*BC* terminating queue. This is possible because the existing implementation of the terminating queue used a modularized queue that can be easily replaced.

Dequeue operation can be non-deterministic in this queue. For example, the state QAsB0C represents a queue with zero or more A followed by a B. A queue in state QAsB0C accepts an dequeue of A and an dequeue of B. The state of the A*BC* queue is further split into empty states and non-empty states. For example, the state Q0C is an empty state and rejects all dequeue operations. As shown in Figure 36, the implementation of the A*BC* queue is written in the same interface as the capacity-limited queue.

When the divergence-free test, trace refinement test and failure refinement test are applied to a system using the A*BC* queue, the test result shows the same pattern as the earlier test on capacity limited queue.

```
datatype TypeQueue= Q0A | QAAs | QAsB0C | QAsBCCs | Q0C  | QCCs
qEmpty(Q0A)=True
qEmpty(Q0C)=True
qEmpty(_)=False

qNewQueue=Q0A

qEnqueue(Q0A,    A)=QAAs
qEnqueue(Q0A,    B)=QAsB0C
qEnqueue(QAAs,   A)=QAAs
qEnqueue(QAAs,   B)=QAsB0C
qEnqueue(QAsB0C, C)=QAsB0C
qEnqueue(QAsBCCs,C)=QAsB0C
qEnqueue(Q0C,    C)=QCCs
qEnqueue(QCCs,   C)=QCCs

qValidEnqueue(Q0A)={A,B}
qValidEnqueue(QAAs)={A,B}
qValidEnqueue(_)={C}

qCanEnqueue(x)={q | q←TypeQueue, member(x,qValidEnqueue(q))}

qDequeue(QAAs)=  {(A,Q0A),(A,QAAs)}
qDequeue(QAsB0C)={(A,QAsB0C),(B,Q0C)}
qDequeue(QAsBCCs)={(A,QAsBCCs),(B,QCCs)}
qDequeue(QCCs)={(C,Q0C),(C,QCCs)}
qDequeue(_)={}
```

Figure 36: CSP Implementation of capacity-limited queue

6 Other objects

There are several synchronization objects we tested in the project, but testing on these object are not presented in this project due to the word limit. In this section, we briefly describe their interface. Interested readers can refer to Anonymous Github Repository at `https://anonymous.4open.science/r/Oxford-Year3Project-8DEC`

Barrier is another synchronization object commonly used in concurrent programming. A sync call returns only after all processes using the barrier calls sync.

```
trait BarrierT{
  def sync: Unit
}
```

With the exchanger object, processes can exchange data with another process. Like the MenWomen object, processes pair up and return the data of the other process.

```
trait ExchangerT[A]{
  /** Exchange x with another thread. */
  def exchange(x: A): A
}
```

Filterchan is like a channel but a process can specify what data it wants to receive by providing a function as a parameter in receive.

```
trait FilterChanT[A]{
  /** Send x on the channel. */
  def send(x: A): Unit

  /** Receive a value that satisfies p. */
  def receive(p: A => Boolean): A
}
```

7 Conclusion

The synchronization test technique can be used to test the correctness of a synchronization object given its interface and expected behaviour. In the thesis, we describe and apply the linearization test to many synchronization objects, and these tests can distinguish between the correct and faulty implementation of the object. We show how to optimize the linearization test by reducing the redundancy in the linearizer process and compare the optimized linearization test with the explicit linearization point test. With all the techniques, we apply the linearization test to a complicated object.

The linearization test has some drawbacks. The complexity of the testing system usually grows exponentially regarding the number of processes and the size of the variable used. As a result, we usually use a small set of processes and variable values, which usually suffices to find a counterexample trace in a faulty system. Besides, sometimes reducing the variable size has no side effects. In the MenWomen object section, the variable `stage` is declared as a Scala integer, but only value 0,1,2 is used.

Also, some of the synchronization objects tested in the thesis are more like artificially created objects for concurrent teaching.

For future work, one can build CSP models for synchronization objects from other sources, such as JVM Source code. When more CSP modules become available, it will also be interesting to build a compiler from Java to CSP to enable large scale automatic testing of synchronization objects. Finally, one can look into improving linearization testing further to allow efficient testing of large systems.

8 Reference

References

[GRABR14] Thomas Gibson-Robinson, Philip Armstrong, Alexandre Boulgakov, and A.W. Roscoe. FDR3 — A Modern Refinement Checker for CSP. In Erika Ábrahám and Klaus Havelund, editors, *Tools and Algorithms for the Construction and Analysis of Systems*, volume 8413 of *Lecture Notes in Computer Science*, pages 187–201, 2014.

[GRPAR13] Thomas Gibson-Robinson, Alexandre Boulgakov Philip Armstrong, and A.W. Roscoe. *Failures Divergences Refinement (FDR) Version 3*, 2013.

[Lea06] Douglas Lea. *Concurrent Programming in Java(TM): Design Principles and Patterns (3rd Edition) (Java (Addison-Wesley))*. Addison-Wesley Professional, 2006.

[Low17] Gavin Lowe. Testing for linearizability. *Concurrency and Computation: Practice and Experience*, 29(4), 2017.

[Low21] Gavin Lowe. Synchronisationlinearisation. `https://github.com/GavinLowe1967/SynchronisationLinearisation`, 2021.

[Low22] Gavin Lowe. Parameterized verification of systems with component identities, using view abstraction. *International Journal on Software Tools for Technology Transfer*, 24(2):287–324, Apr 2022.

[Ros10] A.W. Roscoe. *Understanding Concurrent Systems*. Springer-Verlag, Berlin, Heidelberg, 1st edition, 2010.

[SI09] Konstantin Serebryany and Timur Iskhodzhanov. Threadsanitizer: data race detection in practice. In *Proceedings of the workshop on binary instrumentation and applications*, pages 62–71, 2009.