1 Introduction

A concurrent datatype is ... A concurrent datatype offers encapsulation of concurrency and makes writing concurrent programs simpler.

For example, the MenWomen object is a concurrent datatype that captures a classical synchronization problem. In this problem, some processes need to pair with other processes by exchanging their identities. Figure 1 is an interface of a concurrent datatype for the MenWomen problem.

```
trait MenWomenT{
  def manSync(me: Int): Int
  def womanSync(me: Int): Int
}
```

Figure 1: Interface

The safety and liveness properties are essential to the correctness of concurrent datatypes. The safety property states that the behaviour of the concurrent object should observe some invariant. For example, if a process with identity 1 calling manSync returns 2, then the process with identity 2 should call womanSync and return 1. The liveness property states that the concurrent object should not block all synchronization when synchronization is possible between one or more processes. For example, a system with one process calling manSync and one process calling womanSync should not deadlock.

In this paper, we examine the above two correctness properties for various concurrent datatypes. In addition, we provide a few CSP implementations for objects commonly used in concurrent programming, which can be used in future CSP projects.

1.1 Linearization test

To verify the correctness of a concurrent datatype, one can carry out the Linearization test described in [TODO: Reference]. The linearization testing framework measures each call's starting and returning time to get a history of function calls and function returns. Then for the observed history, the testing framework attempt to find a series of synchronization point that obeys the safety property. If the framework can not find a valid synchronization point series, then the concurrent datatype implementation is wrong.

In this remaining section we shall look at a few history from the MenWomen object. The timeline in Figure 2 visualizes the function call history of two process T1 and T2. T1 first calls manSync, then T2 calls womanSync. A synchronization occurs between T1 and T2. T1 returns the identity of process T2 then T2 returns the identity of process T1.

manSync(T1):T2

T1   Call.T1.manSync   Return.T1.manSync.T2

Sync.T1.T2.T2.T1

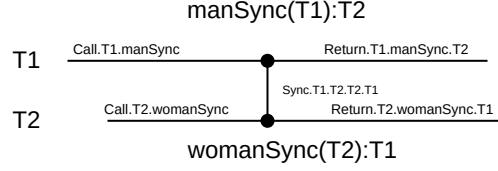T2   Call.T2.womanSync   Return.T2.womanSync.T1

womanSync(T2):T1

Figure 2: Visualized history of T1 calling manSync and T2 calling womanSync

In Figure 3, both processes calls manSync, and no synchronization is possible. Note that the liveness condition is not invalidated even if the system deadlocks in this case.

manSync(T1)

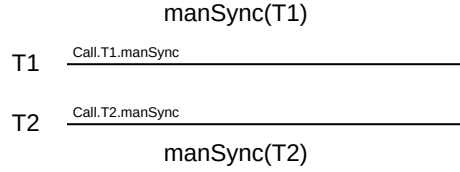T1   Call.T1.manSync

T2   Call.T2.manSync

manSync(T2)

Figure 3: Visualized history of both T1 and T2 calling womanSync

Scheduling is one of the reasons validating a history can be complicated. In Figure 4, process T3 calls manSync first but gets descheduled. Then T1 calls manSync and synchronizes with T2 which later calls womanSync. To find a valid series of synchronization point, the linearization framework usually needs to search a large state.
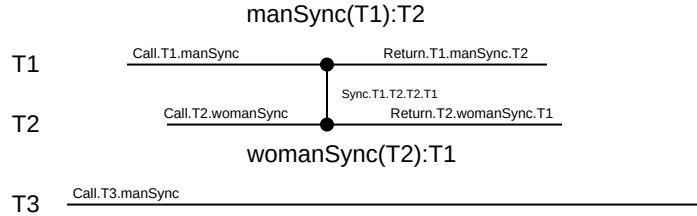
manSync(T1):T2

T1   Call.T1.manSync   Return.T1.manSync.T2

Sync.T1.T2.T2.T1

T2   Call.T2.womanSync   Return.T2.womanSync.T1

womanSync(T2):T1

T3   Call.T3.manSync

Figure 4: Visualized history of T3 get descheduled

1.2 Checking safety property using CSP

The history can be captured as a trace of a system. In addition to performing the function body, each function call sends a Call event before the function body and a Return event after the function body. Figure 5 is the definition of Call and Return channel in CSP. The definition of the channel usually need to include the identity of the calling process, the function called, and its parameter.

```
——identity of the calling process
——function called by the process
channel Call : TypeThreadID . TypeOps
——identity of the calling process
——function called by the process
——return value of the function call
channel Return: TypeThreadID . TypeOps . TypeThreadID
```

Figure 5: Definition of call

To check the safety property, we check that a testing system built from some processes using the concurrent datatype refines a specification process built from the object definition in CSP trace model.

A generic and scalable system is used for the testing system to generate possible histories of processes using concrete datatype. Each process in the testing system can call any function from the concurrent object with any arguments allowed. Such construction allows to build a system with an arbitrary number of processes, and cover as many histories as possible. The specification generates all valid histories. The process uses the same number of linearizer processes synchronizing on events from the Sync channel. Each linearizer process repeatedly calls a function, synchronize with zero or some processes, and returns according to the calling argument and extra information from the synchronization point. We shall see a concrete testing system and specification process in the MenWomen section.

1.3 Checking liveness property using CSP

For liveness property, we check the same generic and scalable testing system refines the same specification process, but in the failure model. Suppose all process calls manSync. Since a linearizer process calling manSync sends Return event only after synchronizing Sync event with another linearizer process calling womanSync, the linearizer will refuse to return any function call, which is a expected behavior. One can use a datatype-specific specification process that does not explicitly use any synchronization points. However, reusing the linearizer process is easier.

2 Common Objects

2.1 Shared Variable

The usage of shared variables is common in concurrent datatypes. For example, some concurrent datatypes may temporarily store the identity of a waiting process. However, CSP is more like a functional programming language and does not support mutable variables.

A recursive process in CSP can capture the behaviour of a shared variable. The recursive process holds the value of the variable in its parameter. At any time, the variable process is willing to answer a query for the variable value in channel getValue. Alternatively, the process can receive an update on the variable value in channel getValue, after which the function recurses with the new variable value.

Because it is natural for a concurrent datatype to use multiple shared variables, the global variable is implemented as a CSP module, as shown in Figure 6. The module requires two parameters. TypeValue is the set of possible values for the variable, and initialValue is the value before any process modifies the variable. An uninitialized variable module is also available in the same figure 6, with the only difference that the variable non-deterministically chooses an initial value from TypeValue at start time. runWith is a convenient helper function to run a given process P with the Var process. If the parameter hide is true, runWith function hides all events introduced by the shared variable. In later chapters, we will see how the runWith function helps reduce the code complexity of the synchronization object implementation.

```
module ModuleVariable(TypeValue, initialValue)
  Var(value) = getValue!value → Var(value)
             □ setValue?value → Var(value)
  chanset = {|getValue, setValue|}
exports
  //(Bool, Proc) → Proc
  runWith(hide,P) = if hide then (Var( initialValue ) [|chanset|] P) \ chanset
                            else   Var( initialValue ) [|chanset|] P
  channel getValue, setValue: TypeValue
endmodule


module ModuleUninitVariable(TypeValue)
  Var(value) = getValue!value → Var(value)
             □ setValue?value → Var(value)
  chanset = {|getValue, setValue|}
exports
  runWith(hide,P) =
    if  hide then ((|~| x:TypeValue • Var(x)) [| chanset |] P) \ chanset
    else  (|~| x:TypeValue • Var(x)) [| chanset |] P
  channel getValue, setValue: TypeValue
endmodule
```

Figure 6: CSP implementation of global variable process module

Figure 7 is an example of two processes using a shared variable. The first line in the example creates a shared variable VarA with value ranging from 0 to 2 and initialized with 0. Process P increments VarA modulo 3 forever and process

Q reads VarA forever. Process P interleaves with process Q, and the combined process is further synchronized with the variable VarA process. In the resulting process System, changes to VarA made by process P is visible to process Q.

```
instance VarA = ModuleVariable({0..2},0)
P = VarA::getValue?a → VarA::setValue!((a+1)%3) → P
Q = VarA::getValue?a → Q
System = VarA::runWith(false,P|||Q)
```

Figure 7: CSP Example of a process using a shared variable

2.2 Semaphore

Semaphore is a simple but powerful concurrent primitive. The essay shall describe and use a simplified binary semaphore from [TODO: Reference], which removes interrupts and timeout operations.

A binary semaphore can either be raised or unraised. A down function call, usually by a process that previously acquired the semaphore, raises the semaphore regardless of the semaphore state. If a process calls the up method when the semaphore is raised, the semaphore becomes unraised. However, if the semaphore is unraised, the process waits until another process calls down and proceeds to put down the semaphore. Depending on the initial state of the semaphore, a binary semaphore can be further classified as a mutex semaphore or a signalling semaphore.

Modelling a semaphore is simple in CSP. A process may call up function or down function via channel upChan or channel downChan respectively. The semaphore is modelled by a process implemented by two mutually recursive functions Semaphore(True) and Semaphore(False). The semaphore process representing an unraised state accepts a upChan event by any process and proceeds to the raised process. The semaphore process representing a raised state can either accept a upChan event and recurse to the raised process, or accept a downChan event and proceed to the unraised process.

2.3 Monitor

A Monitor is another powerful concurrent primitive. This essay will also use a simplified monitor from [TODO:reference].

Usually, code (protected by the monitor) is wrapped inside a synchronized block, as shown in Figure 9. But before the process runs code inside the synchronized block, the process needs to check if others have obtained the monitor lock. Also, an exiting process needs to release the monitor lock so that other processes can continue. In Figure 10, edge 1 represents a process entering the synchronized block, and edge 6 describes the transition of a process exiting the synchronized block. Inside the synchronized block, the process may call notify,

```
module ModuleSemaphore(TypeThreadID, initialState)
  −−Raised
  Semaphore(True) = downChan?id → Semaphore(False)
                  □ upChan?id → Semaphore(True)
  −−Unraised
  Semaphore(False)= upChan?id   → Semaphore(True)

  chanset = {|upChan, downChan|}
exports
  −−runWith::(Bool,Proc) → Proc
  runWith(hide,P) = (Mutex [| chanset |] P)  \
                    (if  hide  then chanset else {})
  channel upChan, downChan: TypeThreadID
endmodule
```

Figure 8: Implementation of a binary semaphore in CSP

```
class MenWomen {
  private var a = 0;
  def op1():Unit = synchronized{
    a+=1;
  }
}
```

Figure 9: Description

notifyAll, wait function. A running process calling wait blocks execution and waits
to be notified. When a running process notifies a waiting process, the waiting
process becomes runnable and waits to obtain the monitor lock. Note that the
waiting process can only continue execution after the current process exits its
synchronized block, unlike that a process waiting to acquire the semaphore can
immediately continue after another process raises the semaphore. All above
transitions are represented by edge 2,3,4,5 in diagram 10 respectively.

In the CSP implementation of monitors, the monitor process does most
of the work. Most monitor methods are implemented simply by sending an
event to the monitor process, except for the wait. Recall that a running process
calling wait first waits for notification and reobtains the monitor lock. So after
sending the Wait event, the process waits for a notify event or a spurious wakeup
event. Finally, the waiting process remains for a WaitEnter event, which gives
the process the monitor lock. In practice, the wait function is often guarded with
a while loop to prevent spurious wakeup. Again CSPM does not have a built-in
while loop, and a generic while statement is implemented in continuation pass
style. (Maybe leave this continuation pass style while loop in appendix?)

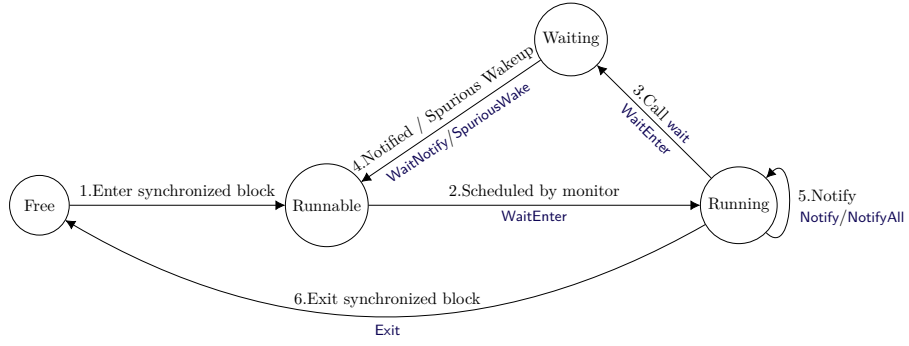The monitor process has two states. When there is no running process,

Figure 10: Time Line

the monitor process can allow a free process to enter its synchronized block by synchronizing on a waitEnter event with the process. When there is a running process, the server process should respond to method calls from the running process, but not allow a process to obtain the monitor lock and proceed to its synchronized block. In either state, the monitor process can spuriously wake up a waiting process.

There is a corner case to highlight in the CSP implementation. In Java, if a running process calls notify method when there is no waiting processing, the function call does nothing. To make sure that the running process is not blocked in this case, when the running process calls notify method, the process send a notify to the server. Then if there are process waiting to be notified, the server process sends a waitNotify event to the waiting process.

```
module ModuleMonitor(TypeThreadID)
channel
  Notify, NotifyAll, Exit, Wait,
  WaitNotify, WaitEnter, SpuriousWake: TypeThreadID

chanset = {| Notify, NotifyAll, Exit, Wait, WaitNotify, WaitEnter, SpuriousWake|}

−−A list of event for every event e in s
repeat(ch, s) =
  if s = {} then SKIP
  else ch?a:s → repeat(ch, diff(s, {a}))

−−cur is current active running thread
−−waiting is a set of threads waiting to be notified
active(cur, waiting) =
  −−current running thread notify
  Notify.cur → (
    if waiting = {} then
      −−do nothing if no thread is waiting
      active(cur, {})
    else
      −−wakeup a process
      WaitNotify?a:waiting →
      active(cur, diff(waiting, {a}))
  ) □
  −−current running thread notifyAll
  NotifyAll.cur → (
    repeat(WaitNotify, waiting);
    active(cur, {})
  ) □
  −−current running thread exit
  Exit.cur → (
    inactive(waiting)
  ) □
  −−current running thread wait
  Wait.cur → (
    inactive(union(waiting,{cur}))
  ) □
  −−spurious wakeup
  waiting ≠ {} & SpuriousWake?a:waiting → (
    active(cur, diff(waiting, {a}))
  )

−−when no active thread is running
```

```
  inactive (waiting) =
    −−pick a thread that is ready to enter
    WaitEnter?a → (
      active(a, waiting)
    ) □
    −−spurious wakeup
    waiting ≠ {} & SpuriousWake?a:waiting → (
      inactive ( diff (waiting, {a}))
    )

exports
  −−Given a process that uses the monitor
  −−Return the process synchronized with the monitor server process
  −−If hide is true, monitor channels are hidden
  runWith(hideSpurious, hideInternal, P) =
    let hideset0 = if hideInternal then chanset else {} within
    let hideset1 = if hideSpurious then hideset0 else diff (hideset0,{|SpuriousWake|}) within
    ( inactive ({}) [|chanset|] P) \ hideset1

  −−java−like synchronized function
  synchronized(me, P)=
    WaitEnter.me →
    P;
    Exit.me →
    SKIP

  enter(me) =
    WaitEnter.me →
    SKIP

  exit (me) =
    Exit.me →
    SKIP

  −−notify()
  notify (me) =
    Notify.me →
    SKIP

  −−notifyAll()
  notifyAll (me) =
    NotifyAll.me →
    SKIP

  −−wait()
```

```
  wait(me) =
    Wait . me → ((
        WaitNotify . me →
        WaitEnter . me →
        SKIP
      ) □ (
        SpuriousWake . me →
        WaitEnter . me →
        SKIP
      )
    )

  whileWait(me,cond) =
    while(cond)(wait(me);SKIP)
endmodule
```

Figure 11: Implementation of Monitor in CSP

3   MenWomen

For simplicity, a process calling ManSync is called a man process, and a process calling WomanSync is called a woman process.

3.1   Implementation

One way to implement the MenWomen object is to use a monitor and a shared variable indicating the stage of synchronization. Figure 12 is a Scala implementation of the MenWomen object with monitor.

- A man process enters the synchronization and waits until the current stage is 0. Then in stage 0, the man process sets the global variable him inside the MenWomen object to its identity. Then the man process notifies all processes so that a waiting woman process can continue. Finally, the man process waits for stage 2.

- A women process enters the synchronization and waits until the current stage is 1. The woman process sets the global variable her to its identity and returns the value of global variable him.

- In stage 2, the waiting man process in stage 0 is wakened up by the woman process in stage 1. The man process notifies all waiting processes and returns the value of her.

The code snippet in Figure 12 is a Scala implementation of the MenWomen process using a monitor by Gavin Lowe. The Scala code is further translated to a CSP code in Figure 13. Every function call begins with a Call event containing

all parameters. And every function call ends with a Return event containing the return value.

```
class MenWomen extends MenWomenT{
  private var stage = 0
  private var him = −1
  private var her = −1

  def manSync(me: Int): Int = synchronized{
    while(stage != 0) wait()
    him = me; stage = 1; notifyAll()
    while(stage != 2) wait
    stage = 0; notifyAll (); her
  }

  def womanSync(me: Int): Int = synchronized{
    while(stage != 1) wait
    her = me; stage = 2; notifyAll();
  }
}
```

Figure 12: Scala implementation of the MenWomen process using a monitor

```
instance  VarStage  =  ModuleVariable({0,1,2},0)
instance  VarHim  =  ModuleUninitVariable(TypeThreadID)
instance  VarHer  =  ModuleUninitVariable(TypeThreadID)
instance  Monitor  =  ModuleMonitor(TypeThreadID)
manSync(me)  =
  Call ! me ! ManSync →
  Monitor:: enter(me);
    Monitor::whileWait(me,  \ ktrue,kfalse •
      VarStage::getValue ? x →
      if  x≠0 then ktrue else  kfalse
    );
    VarHim::setValue ! me →
    VarStage::setValue ! 1 →
    Monitor:: notifyAll (me);
    Monitor::whileWait(me,  \ ktrue,kfalse •
      VarStage::getValue ? x →
      if  x≠2 then ktrue else  kfalse
    );
    VarStage::setValue ! 0 →
    Monitor:: notifyAll (me);
    VarHer::getValue ? ans →(
  Monitor:: exit (me);
  Return ! me ! ManSync ! ans→
  SKIP
  )
womanSync(me)=
  Call ! me ! WomanSync →
  Monitor:: enter(me);
    Monitor::whileWait(me,  \ ktrue,kfalse •
      VarStage::getValue ? x →
      if  x≠1 then ktrue else  kfalse
    );
    VarHer::setValue ! me →
    VarStage::setValue ! 2 →
    Monitor:: notifyAll (me);
    VarHim::getValue ? ans →(
  Monitor:: exit (me);
  Return ! me ! WomanSync ! ans→
  SKIP
  )
```

Figure 13: CSP implementation of the MenWomen

3.2 Linearization Test

//This technique mostly applies to all objects.

A generic and scalable testing system is introduced to make the test cover as many cases as possible. The testing system includes an arbitrary number of processes, and each process calls any functions provided by the synchronization objects for any finite or infinite number of times. For the CSP implementation, each process non deterministically chooses a function and its parameters, or the process chooses to terminate. Figure 15 is the definition of a testing process and a testing system for MenWomen object in CSP.

```
Thread(me)=
    (manSync(me);Thread(me))
 ⊓(womanSync(me);Thread(me))
 ⊓STOP
System(All)=runWith(True,True,||| me:All • Thread(me))
```

Figure 14: CSP implementation of the testing processes and system

There are two properties one should check for the system. First, the traces of the system are valid. If a thread t_1 calls ManSync and returns the thread identity t_2, then t_2 must call WomanSync somewhere. This is done by checking that the system refines a linearisation specification (introduced earlier?). Second, the system should not deadlock when matching is possible. This can be done by another specification process, which records the set of men processes and women processes to determine if matching is possible. However, it is easier to modify the linearisation specification process and check that the system failure refines the specification process.

The implementation of the linearizer specification relies on the synchronization point. The synchronization point of a man process and a woman process is represented by a Sync event. The $Sync.t_1.a.t_2.b$ represents the synchronization of a man process t_1 with parameter a with a woman process t_2 with parameter b.

With the Sync event, it is simple to define the lineariser for a process. The process can choose to call ManSync, synchronize with a woman process, and return the identity of the woman process. Or, the process can choose to call WomanSync, synchronize with a man process, and return the identity of man process. Also, the process chooses to terminate when it finishes calling a function.

4 ABC

In the ABC object, three threads are involved in each round of synchronization. For simplicity, a process calling syncA, syncB, and syncC is called a A process, B process, C process. In each round of synchronization, A A-process, B-process,

```
Lin( All ,me)= (
    Call ! me ! ManSync→
    Sync ! me ? mereturn ? other ? otherreturn →
    Return ! me ! ManSync ! mereturn →
    Lin( All ,me)
)⊓(
    Call ! me ! WomanSync →
    Sync ? other ? otherreturn ! me ? mereturn →
    Return ! me ! WomanSync ! mereturn →
    Lin( All ,me)
)⊓STOP
```

Figure 15: CSP implementation of the testing processes and system

and C-process synchronize, and each process returns the argument of two other processes.

```scala
class ABC[A,B,C] extends ABCT[A,B,C]{
  // The identities of the current (or previous) threads.
  private var a: A = _
  private var b: B = _
  private var c: C = _

  // Semaphores to signal that threads can write their identities.
  private val aClear = MutexSemaphore()
  private val bClear, cClear = SignallingSemaphore()

  // Semaphores to signal that threads can collect their results.
  private val aSignal, bSignal, cSignal = SignallingSemaphore()

  def syncA(me: A) = {
    aClear.down          // (A1)
    a = me; bClear.up    // signal to b at (B1)
    aSignal.down         // (A2)
    val result = (b,c)
    bSignal.up           // signal to b at (B2)
    result
  }

  def syncB(me: B) = {
    bClear.down          // (B1)
    b = me; cClear.up    // signal to C at (C1)
    bSignal.down         // (B2)
    val result = (a,c)
    cSignal.up           // signal to c at (C2)
    result
  }

  def syncC(me: C) = {
    cClear.down          // (C1)
    c = me; aSignal.up   // signal to A at (A2)
    cSignal.down         // (C2)
    val result = (a,b)
    aClear.up            // signal to an A on the next round at (A1)
    result
  }
}
```

Figure 16: Scala implementation of the ABC using semaphores

For the above semaphore implementation of ABC object, In each round

- Initially semaphore aClear is raised.

16

- An A-process acquire semaphore aClear, sets the shared variable a to its parameter, raises semaphore bClear and waits to acquire semaphore aSignal. A B-process and a C-process operates in turn with a slight change of semaphore and variable name.

- The A-process is able to continue after a C-process raises semaphore aSignal. The A-process reads the shared variable b and c, raises the semaphore bSignal, and returns. This also happens in turn for the B-process and the C-process.

Using the shared variable and semaphore module, it is easy to translate the Scala implementation to a CSP implementation.

Unlike monitor in Java and Scala, raising a semaphore immediately allows another thread waiting to acquire the semaphore to continue. So in the semaphore implementation, it is essential to take a copy of the two other arguments before raising the semaphore.

On the other hand, what if the implementation of syncA does not take a copy of the argument? It turns out the ABC object still works correctly when only three threads are involved, but fails the linearisation test with four threads.

4.1   Testing

For the MenWomen object, Using the standard linearization testing technique, the following Sync channel can be used to represent the synchronization of three involved threads. For example, the event $Sync.t_1.a.b.c.t_2.d.e.f.t_3.g.h.i$ represents the synchronizations of three threads, t_1, t_2, t_3, in which the first process t_1 calls aSync with a and returns (b, c), the second process t_2 calls bSync with d and returns (e, f), and last process t_3 calls cSync with g and returns (h, i). The spec process should then check that for each synchronization point, the return value of each functional call is the pair of arguments of the two other function call.

```
channel Sync: TypeThreadID . TypeData . TypeData . TypeData .
              TypeThreadID . TypeData . TypeData . TypeData .
              TypeThreadID . TypeData . TypeData . TypeData

Spec = Sync ? aid ? a ? b ? c
            ? bid: diff (TypeThreadID,{aid}) ! b ! a ! c
            ? cid: diff (TypeThreadID,{aid,bid}) ! c ! a ! b
      → Spec
```

//Preparation of test case: I will describe There are two test cases. The first test case involves three threads. Each of the thread chooses a data non-deterministically and then calls one of aSync, bSync, cSync. The second test case involves four thread, which chooses a data non-deterministically and calls aSync. In both cases, the systems should be traced refined (is it this direction in words)

by the specification process. In addition, both systems should never deadlock. Because in both system, there are always threads willing to communicate as aSync, bSync, cSync respectively.

When testing with both the correct and the faulty versions of ABC object, FDR finishes the first test case relatively quickly, but requires a long time to finish the second test. With logging message from FDR, it was found the compilation of specification process took the longest time. //TODO: Table

4.1.1 Speeding up model compilation

Consider the specification process. Let N be the number of threads in the system, M be the size of the set of all possible arguments. The specification process is the alphabetized parallel of N individual linearisers. In each lineariser, the process first chooses to perform one of aSync, bSync or cSync, chooses the argument of the functional call for Call event, then chooses the rest of arguments for Sync event, and finally performs one event before recursing into itself.

There are $O(3 * N^3 M^9)$ different transitions before the individual lineariser recurses into itself. However, according to the specification process, once the argument and return value of syncA is determined, all remaining arguments and return value are also determined. So only $O(3 * N^3 M^3)$ transitions are valid according to the specification.

```
Call ! me ! ASync ? a →
  Sync ! me ! a ? b ? c
       ? t2: diff ( All ,{me}) ? t2b ? t2a ? t2c
       ? t3: diff ( All ,{me,t2}) ? t3c ? t3a ? t3b →
  Return ! me ! ASync ! b ! c →
  Lin( All ,me)
```

With the above analysis, it is tempting to optimize the individual lineariser by using the information from the specification process. Instead of choosing all possible remaining arguments, the individual lineariser could choose arguments that are correct according to the specification process.

```
Sync ! me ! a ? b ? c
     ? t2: diff ( All ,{me}) ! b ! a ! c
     ? t3: diff ( All ,{me,t2}) ! c ! a ! b →
Return ! me ! ASync ! b ! c
```

It is possible to further simplify the Sync channel, as now the arguments representing return value are redundant. This change does not reduce the number of transitions for an individual lineariser, but it may help FDR simulating the model faster.

With the above optimizations, the testing finishes quickly for both test cases.

```
channel Sync: TypeThreadID . TypeThreadID . TypeThreadID .
             TypeData . TypeData . TypeData

Lin( All ,me)= (
   Call ! me! ASync ? a →
   Sync ! me ? t2:diff( All ,{me}) ? t3:diff( All ,{me,t2}) ! a ? b ? c →
   Return ! me! ASync! b! c →
   Lin( All ,me)
) . . .
```

//TOADD: Table of compilation time

4.1.2   Explanation of the error case

With the traces of the counterexample from FDR, it is possible to see what goes wrong in the faulty version when there are four threads.
   //TODO: Draw diagram using Scala code for this.

- Thread T_A, T_B, T_C call aSync, bSync or cSync respectively, and put down its argument in turn.

- Thread T_A raises bSignal without saving a copy of return value (b,c). The other two threads T_B, T_C are able to continue and exit.

- Thread T_D, T_B, T_C call aSync, bSync or cSync respectively, and put down its argument in turn.

- Thread T_A uses the wrong overwritten (b, c) as return value.

4.1.3   Conclusion

In the above section, we tested a semaphore based concurrent datatypes. With the testing result, we showed that it is important to be reminded that a thread waiting for the semaphore to raise can immediately continue and overwrite shared variables, after the semaphore is raised by another process.